

AI Usage Disclosure

The AI Bidding War: Autonomous Procurement & Vendor Negotiation

Asli Gulcur (solo project) · CMU Agentic Systems Studio · Track A · April 2026

This file describes **generative AI and assistant tools** used in the Phase 2 of *The AI Bidding War*, per course expectations: what was used, for what purpose, what was edited manually, and what was **independently verified**.

Human-Led Scope

I **drove** the design, **governance** decisions (Strategic Audit, human gatekeeper, deterministic guardrails, optional nudges), and **architecture** — including choosing **semi-autonomous** operation (orchestrated roles + human approval where appropriate) over full end-to-end autonomy — as appropriate for a **procurement / policy** use case that needs guardrails/governance.

Strategic Audit independence: I decided the audit should **not** interact with the **buyer** in the loop (no buyer–auditor dialogue or automatic injection of the full audit into the buyer’s next turn); the **human** reviews the Strategic Audit report, then **alone** may issue a short **nudge** to the buyer after a rejection — preserving separation between negotiation and assurance.

Importance of Human In The Loop: It was important to me that human was accountable for the final decision, and was also empowered to intervene after reading the Strategic Audit Report during the scenario by nudging/redirecting the Buyer agent.

Evaluation Design: I also decided **what to simulate and how hard to stress it: scenario planning** across `config/scenarios.json` (baseline, compliance, budget pressure, conflict-of-interest, etc.). I wanted to inject **adversarial friction**, and created the **malicious buyer** persona (`agents/malicious_buyer/`) to validate guardrails under pressure — not just happy-path negotiation.

Offline CFO narrative report (`cfo_monthly_report.py` + `agents/cfo_narrative/agent.md`) was my choice for **portfolio-level** reflection on evidence logs, distinct from in-run Strategic Audit. I thought about making it useful for the organization, and how the CFO could benefit from such executive summaries.

Reports, System Diagrams and Evaluation Frameworks: While Cursor provided the initial drafts, I revised and iterated heavily in order to finalize them. The initial reports or the test case summaries were disjointed; lacked strategic narrative or key evaluation filters that would be useful for a human. I drove for clarity, strategic narrative and usability for a human.

The project went through **multiple iterations**: starting from a **narrow** scope and **expanding** deliberately into a **more realistic business simulation** (richer **scenario** set, evaluation pipeline, **CFO-style** reporting), with each step reviewed thoroughly. My focus has also been on the human/user experience and value/impact for the organization.

I independently verified all the scenario runs, the logs and the evaluation artifacts.

Cursor's Role

I used **Cursor** as a **coding agent** in the IDE to help **develop** the project — drafting and refactoring code, suggesting fixes, navigating the repo structure, implementing the solution, and generating the necessary files. It also acted as a **sparring partner**: **brainstorming** design and implementation options with me (not replacing my choices) and **explaining development steps** — for example how pieces fit together, what to run next, and how to debug — so I could **learn the workflow** while remaining responsible for architecture and governance decisions.

I occasionally checked with Anthropic's Opus 4.6 for ad-hoc questions.