

- Canvas submission packet — The AI Bidding War
 - Project title
 - Team members
 - Selected track
 - One-paragraph summary
 - Headline results (evaluation harvest **2026-04-18**)
 - Submission bundle — repository or project-file link
 - Project video (included in the zip)
 - Final report
 - Architecture diagram
 - Screenshots (4–8+ with index)
 - Evaluation summary
 - AI usage disclosure (Phase 3 — PDF only)
 - List of submitted files and folders (inside `Final Project AI Bidding War.zip`)
 - Reviewer map ($\leq$ ~10 minutes)

Canvas submission packet — The AI Bidding War

Course: Agentic Systems Studio · **Track A — Technical Build**

Purpose: Single PDF for Canvas that points reviewers to the zip, reports, diagram, video files, and evaluation. **Executive HTML + PDF (recommended):** from the project root run `python scripts/render_canvas_submission_packet.py` — writes `CANVAS_SUBMISSION_PACKET_executive.html` and `CANVAS_SUBMISSION_PACKET_executive.pdf` in this folder (pandoc + Chrome; `--html-only` skips PDF). Alternatively export this Markdown manually, or paste sections into your template.

Project title

The AI Bidding War — hub-and-spoke multi-agent procurement with a buyer, three isolated vendors (A/B/C), Strategic Audit (LLM), human approval, governance nudges, and a deterministic **\$15,000** award guardrail.

Team members

Name	Role
Asli Gulcur	Solo author — design, implementation, evaluation

Selected track

Track A — Technical Build (CMU Agentic Systems Studio)

LLM / ANTHROPIC_MODEL : Default `claude-sonnet-4-6` (`.env`, orchestrator, advisor, Docker fallbacks, `agent.md`). **Anthropic** has announced deprecation of `claude-sonnet-4-20250514` (Phase 2 Model) in **June 2026**; archived Phase 2 logs use that id.

One-paragraph summary

Enterprise-style procurement negotiation is structured, auditable, and failure-prone (budget overruns, steering, conflict of interest). **The AI Bidding War** is a runnable **Python CLI** that simulates multi-round RFP negotiation: isolated vendor threads, a buyer hub, an LLM **Strategic Audit**, **human approval** gates, optional governance nudges, and a **code-enforced** `award_contract` guardrail. Evaluation is harvested from `logs/` into `evaluation/` CSVs and a generated final report. The design keeps final authority out of the model for awards and documents limitations (e.g. prompt-level public buyer channel).

Headline results (evaluation harvest 2026-04-18)

From `evaluation/test_cases.csv` (pinned in `evaluation/version_notes.md`):

- **11** scenario runs harvested (scenarios **1–11**), exceeding the syllabus minimum of **5** completed test scenarios.
- **Deal outcomes:** **6** runs with `Status = SUCCESS` (in-budget awards where applicable); **5** runs with `Status = NO_AWARD` (human veto / governance paths — scenarios **2, 6, 7, 10, 11** for this harvest).
- **Deterministic guardrail:** No award executes **above** the **\$15,000** hard ceiling in code (`skills/award_contract.py`); malicious scenarios are exercised without illegal spend.
- **Failure / boundary documentation:** `evaluation/failure_log.md` records themes **FL-001** through **FL-004** (≥ 2 concrete boundary cases required by the syllabus — here, four tagged themes with scenario cross-references).

Submission bundle — repository or project-file link

Item	Value
Submitted archive	<code>Final Project AI Bidding War.zip</code> — contains the full project folder (see List of submitted files and folders below).

Note (project-file “link”): The zip file `Final Project AI Bidding War.zip` is uploaded to Canvas alongside this PDF. No external repository URL is provided; the **full project** is the contents of that uploaded archive. That is the **project-file link** for this submission: the grader opens the **same Canvas submission** (PDF + zip attachment), not a separate website.

Primary entry inside the zip: project root `README.md` (setup, run: `python orchestrator.py --id <N>`, folder guide, evaluation semantics).

Project video (included in the zip)

Item	Path inside zip
Project video	<code>media/The AI Bidding War - Procurement (5 mins).mov</code>

Model id (video only): Video (e.g. **Scenario 2** — malicious / guardrail demo) was recorded with `ANTHROPIC_MODEL=claude-sonnet-4-20250514` (Phase 2 Model). Anthropic has announced deprecation of that id in **June 2026**; use it **for this demo segment** only, so the **corrupt buyer** arc is easy to see on screen; the submitted project still defaults to `claude-sonnet-4-6` in `.env`. The **2026-04-18** evaluation harvest and `evaluation/failure_log.md` (FL-001) describe **4.6** behavior and **method variance**, not a guardrail change. **Restore** `claude-sonnet-4-6` **after** any session that used `claude-sonnet-4-20250514` (Phase 2 Model).

Final report

Deliverable	Path in zip
Phase 3 final report (PDF)	docs/Phase 3/Phase3_Final_Report_4_21_2026.pdf
Log / evidence analysis (PDF)	docs/Phase 3/Phase3_Log_Evidence_Analysis.pdf
Supplementary evaluation report (PDF)	docs/Phase 3/Phase3_Supplementary_Evaluation_Report.pdf
Phase 2 vs Phase 3 failure themes (FL-001–FL-004, instructor feedback)	evaluation/failure_log.md ; final report §6–8 in docs/Phase 3/Phase3_Final_Report_4_21_2026.pdf ; evaluation/FINAL_EVALUATION_REPORT.md (Phase 2 vs Phase 3 is not a separate memo)
Multi-seed variance (PDF)	docs/Phase 3/Multi Seed Variance Evaluation 4_21_2026.pdf ; see also evaluation/multi_seed_variance/ — runs.csv , MULTI_SEED_VARIANCE_REPORT.md

Architecture diagram

Item	Path
System architecture (SVG)	docs/System Architecture Diagram/ai_bidding_war_system_architecture.svg
Interactive canvases (HTML)	docs/Canvases/ (e.g. MA_ArchCanvas_The_AI_Bidding_War.html , MA_SpecCanvas_The_AI_Bidding_War.html)

Screenshots (4–8+ with index)

Item	Path
Index (captions + mapping)	docs/Screenshots/screenshot_index.md
Images	docs/Screenshots/ — 01_home.png ... 09_cfo_narrative.png

Evaluation summary

Item	Path / note
Harvest date (current pin)	2026-04-18 — see <code>evaluation/version_notes.md</code>
Test cases	<code>evaluation/test_cases.csv</code>
Rubric results	<code>evaluation/evaluation_results.csv</code>
Rubric baseline	<code>evaluation/evaluation_rubric.md</code>
Failure / boundary log	<code>evaluation/failure_log.md</code> (FL-001–FL-004)
Compact table	<code>evaluation/FINAL_RESULTS_TABLE.md</code>
Generated report	<code>evaluation/FINAL_EVALUATION_REPORT.md</code> (from <code>python scripts/export_evaluation.py</code>)
Multi-seed variance (Phase 2 feedback)	<code>evaluation/multi_seed_variance/</code> — <code>runs.csv</code> , <code>MULTI_SEED_VARIANCE_REPORT.md</code> , <code>README.md</code> ; regenerate: <code>python scripts/collect_multi_seed_variance.py</code>

Syllabus filenames (under `evaluation/`): `test_cases.csv`, `evaluation_results.csv`, `failure_log.md`, `version_notes.md` — all present in the zip. `multi_seed_variance/` is additional instructor-feedback evidence, not a syllabus-filename requirement.

AI usage disclosure (Phase 3 — PDF only)

Item	Path in zip
Generative-AI / tooling disclosure	<code>docs/Phase 3/Phase3_AI_USAGE.pdf</code>

List of submitted files and folders (inside Final Project AI Bidding War.zip)

Path	Role
README.md	Navigation, setup, run, evaluation semantics; team members
orchestrator.py , agents/ , skills/ , config/scenarios.json	Runtime
evaluation/	CSVs, rubric, reports, failure log, version notes, multi_seed_variance/
logs/	Session logs, evidence JSON, optional contract.txt sample
scripts/	export_evaluation.py and tooling
docs/	Phase 1–3 docs, diagrams, screenshots, Phase 3 PDFs (including Phase3_Final_Report_4_21_2026.pdf , Phase3_AI_USAGE.pdf , Multi Seed Variance Evaluation 4_21_2026.pdf , CANVAS_SUBMISSION_PACKET_executive.pdf , CANVAS_SUBMISSION_PACKET_executive.html , canvas_packet_executive.css)
media/	The AI Bidding War – Procurement (5 mins).mov (project video)
requirements.txt , package.json	Dependencies (Python / optional Node) — venv + CLI is the supported workflow; Docker files exist but were not used for Phase 3
CFO Reports/	Optional CFO narrative Markdown

Reviewer map (≤ ~10 minutes)

1. Unzip **Final Project AI Bidding War.zip** and read **README.md** .
2. Open **docs/Phase 3/Phase3_Final_Report_4_21_2026.pdf** .
3. Open **docs/System Architecture Diagram/ai_bidding_war_system_architecture.svg** .
4. Skim **docs/Screenshots/screenshot_index.md** and sample PNGs.
5. Open **evaluation/FINAL_RESULTS_TABLE.md** or **evaluation/test_cases.csv** ; confirm harvest in **evaluation/version_notes.md** .
6. Play **media/The AI Bidding War – Procurement (5 mins).mov** (syllabus video elements; cumulative structure — stopping early is fine).
7. Read **docs/Phase 3/Phase3_AI_USAGE.pdf** for generative-AI disclosure.